True Optics in 3D Rendering

Christian Oeien

August 19, 2016

Abstract

A 3D renderer of basic geometric shapes of variable optics is presented. The implementation created permits union and intersection of spheres, planes, cylinders and cones (and their spatial inverses) each of parameterized reflection, refraction, absorption and pass-through on semi-transparency. Some optical phenomena are not accounted for by the written program and this is discussed.

1 Introduction

The presented renderer is a ray-tracer, as described by [3]. An alternative approach for rendering 3D is to project each composed surface onto a raster. Surfaces farther from the observer are overlayed by ones closer. This projection technique is used for real-time rendering of moving 3D, because ray-tracing requires much more calculation. When creating a projection renderer two problems arises immediately. The surfaces may intersect, and even tho they do not intersect a naïve test for which surface to render first will fail in many cases. It is shown in [1] that the problems we encounter are not trivial. And after these complex problems are solved, the result is still not realistic. The surfaces must somehow be shaded according to lighting or reflexion. This would be performed for the visible parts of the surfaces. Also, for the realistic rendering of combined surfaces, there would be need for some shading when applying the colors, as described in [2].

For optical realism the ray-tracer technique is chosen. This method arrives at realistic results by itself, and with much simpler algorithms. It must be noted that the CPU time a program will need is bigger than for the projection technique for a picture of comparable quality because of all the iteration and recursion done in ray-tracing. In the technique of ray-tracing one seeks to calculate the color of each resulting pixel in turn, by tracing the origin of the incoming light-ray from the corresponding direction. Projection simulates the effects of light while ray-tracing models each light-ray itself, on its journey from its various sources. The amount of information collected is limited because we may on each combination perform filtering or addition of three components of the simulated light. If our visual ability was different, say detecting all fre-

Figure 1: Tracer recursion

quencies present like our auditory sense does, this approach would not give a satisfying result.

Physically correct prism effect cannot be modeled with a finite set of traced rays. An attempt would be to implement the ray-casting technique from the light-spot sources, which could be limited to one traced ray per source at each absorbing surface. But light dispersion shall also apply to the source light from other objects, which would still require trace from infinite directions. The presented ray-tracer ignores this problem by considering only a single light-ray from any explicit light-source to an objects surface.

2   Material and Methods

The core engine is written in C and the input-processing and initial composition of the runtime scene of geometrical data for all objects in C++.

The PC architecture is an Intel CPU with `mmxext` and `sse4a` found in `/proc/cpuinfo` running a Linux kernel.

3   Results

The source code is found at `HTTP://github.com/biotty/rmg/graphics/ray`. A separate software library for color manipulation was factored out, and is found under `graphics/rgb`.

For each pixel to generate, the program geometrically follows the corresponding direction to the viewer. The intersection of this ray with each object in the composed 3D scene is calculated. The closest point of intersection to the viewer is then considered for reflexion, refraction and absorption at the hit surface. Functions are called to find colors of incoming light that will arrive in the direction of the viewer, at this point on the surface. The normal vector (direction out of the 3D object) is also needed for these calculations, where we recurse the procedure described, as if we were a viewer at this point for light receiving

Figure 2: Set-intersection of shapes

reflection, refraction and absorption per light point-source. There is of-course no recursion for refraction on an opaque object. Absorption from each spot light source is calculated. When no object is hit when tracing an incoming source light-ray for color, we use the direction to calculate the color by one selected function in `sky.c`. Therefore the spot light sources are not required to render color of a composed scene of objects. Please see fig. 1.

While recursing through functions of optics and geometry we operate a set of bits, one per scene object, telling whether the view-location of our currently traced ray is inside of that scene object. Note that a traced ray may happen to start inside of an opaque scene object, because a transparent object may spatially overlap it and its optics be in force at that location in space.

A basic geometric object must provide two functions to our simulation. For a given ray it calculates the distances to the two intersections through it. Then given a point in space it gives a direction which will be the normal on its surface. Note that the fixation of maximum two intersections means we cannot model a torus: For a practical illustration, observe that four surfaces of a donut touches your knifes blade if you cut it fairly. In general we do not permit any concave (or saddle) surface.

The intersection is the only scene object that is composed of other objects. The intersection of shapes is the space that is such that a point is included in all of those shapes. Since the basic objects have no concave surfaces, an intersection of them does not either. A proof of this is trivial, when using the mentioned observation of maximum one segment (two intersections) of any line through such an object: An intersection of two segments can lead to no more than one segment. Please see fig. 2.

Negatives of the basic shapes except plane are supported (negative planes would still allow for only planes). A negation of a shape is modeled by considering that it has the same enter-intersection and exit-intersection with a ray, but the roles of enter and exit are swapped.

The calculations of intersections are complex but as scene objects they have the same programatic interface, and are traversed isomorphically to any scene

3

Figure 3: Overlap of scene objects

object by the core ray-tracer engine. On the other hand, the support for overlapping scene objects adds complexity to the top-level ray-tracer algorithm itself. A sequential order of scene objects is used to decide precedence when traversing the geometrical boundary of a scene object. If inside of a scene object of higher precedence than the traversed surfaces object, the ray shall be traced as if the surface was not optically hit. We need to know which scene objects we are geometrically but not optically inside of. Please see fig. 3. This is accomplished with a stack data-structure where we push or pop respectively when entering or exiting an object through its surface with no optical presence.

The program takes a minimalistic syntax ASCII-encoded input, and a Python module that feeds this to the ray-tracer based on the equivalent composition using Python objects, is provided. This module also arranges for time-parametric compositions so that sequences of images are generated, to form motion-picture.

On my architecture I observed no difference in CPU time for the three types of floating point. The precision of the calculations for a normal scene do not benefit from the bigger types. The data structures holding the geometrical objects in the 3D scene uses the smallest type for spatial coordinates and a compressed form of colors.

4 Discussion

The code that leverage C 11 and does not compile in C 99 is mostly found in for-loop constructs and idioms trivially rewritten for C 99 if needed. With C 11 a cast between two `struct` pointers where first member type is identical, has defined behaviour by the language. This is relevant for the scene object arguments, and we need C 11 for us to be safe, even if the practice of such casting was common in pre-2011 conformant compilers. Another type-problem is the `struct`s for direction and for point. Casting between these types is possible (the value itself, not a pointer) because the compositions are identical. However, it is not guaranteed to result in a no-op runtime. Also, use of different types would let the compiler stop us from using the wrong type. The source code arranges for the type-safe alternative of using two different `struct`s, by

the alteration of one preprocessor `#define`. The type of floating point used generally in the program is also tunable. This is done with `typedef`.

In C we use `void *` for holding the generic object data. This could benefit from the type-system of C++. The advantage of C is compilation-time.

I do not make use of an image-format software library in my program, but merely feed raw color values and facilitate for feeding this to a picture encoder, like the available `pnmtojpeg` and use of `avconv` on a typical Linux-based operating system.

The limitation on permittable shapes was arbitrary and because I started with the simplest shapes without considering other possibilities. I had already implemented the intersection algorithm when I got aware that there are basic geometrical shapes that do not fit, and that my algorithms had this constraint on intersection-pair. The intersection algorithm was already complex, and there was much work involved to make the floating-points comparisons and the recursive splitting of the sections in order to find the remaining intersection of all component objects. Adding to this complexity would further increase CPU time spent on the tracing, as the part spent in this code already showed to be significant.

It must be noted that the physical phenomenon of optical refraction is frequency-dependent, and this is not modeled by the created ray-tracer. This means that prism-like effects will not occur when tracing through a transparent object. If we were to model this, it would not suffice to trace from only one direction of refraction. And any finite number of directions we chose to trace would yield optically incorrect results.

The function that traces for absorption iterates over all light-sources specified in the scene. If the line between the surface-point and a light-spot does anywhere hit a surface of any scene object then no color is applied from this light-source. This means that total shadow is cast even from a transparent object, and this is not physically correct. The problem is of finding what path a light-ray would take through transparent objects to reach a given point. A solution to this problem would be computationally intensive, and when implemented we would still not obtain the prism effect due to the other problem as described previously. I chose to focus on implementing the optically correct simulations and to conserve simplicity.

References

[1] Walter Heger. Vector hidden line removal and fractional quantitative invisibility. In *Tripod.com*, 1992.

[2] Bui Tuong Phong. Illumination for computer generated pictures. In *Communications of ACM 18(6)*, pages 311–317, 1975.

[3] Paul Rademacher. Ray tracing: Graphics for the masses. In *ACM Crossroads (Summer Issue)*, 1997.